

# Complete LangGraph Sequential Workflow Documentation

## Table of Contents

1. [Introduction](#)
  2. [Prerequisites](#)
  3. [Architecture Overview](#)
  4. [Core Concepts](#)
  5. [Code Breakdown](#)
  6. [Execution Flow](#)
  7. [Installation & Setup](#)
  8. [Usage Examples](#)
  9. [Advanced Customization](#)
  10. [Troubleshooting](#)
- 

## Introduction

This documentation covers a **LangGraph Sequential Workflow Agent** that demonstrates how to build production-ready agentic AI systems. The agent:

- Retrieves documents from a knowledge base
- Generates reasoning using Claude/OpenAI
- Produces validated answers
- Tracks workflow progression

## Use Cases:

- Customer support chatbots
- Research assistants
- Question-answering systems
- Document analysis pipelines
- Automated report generation

## Prerequisites

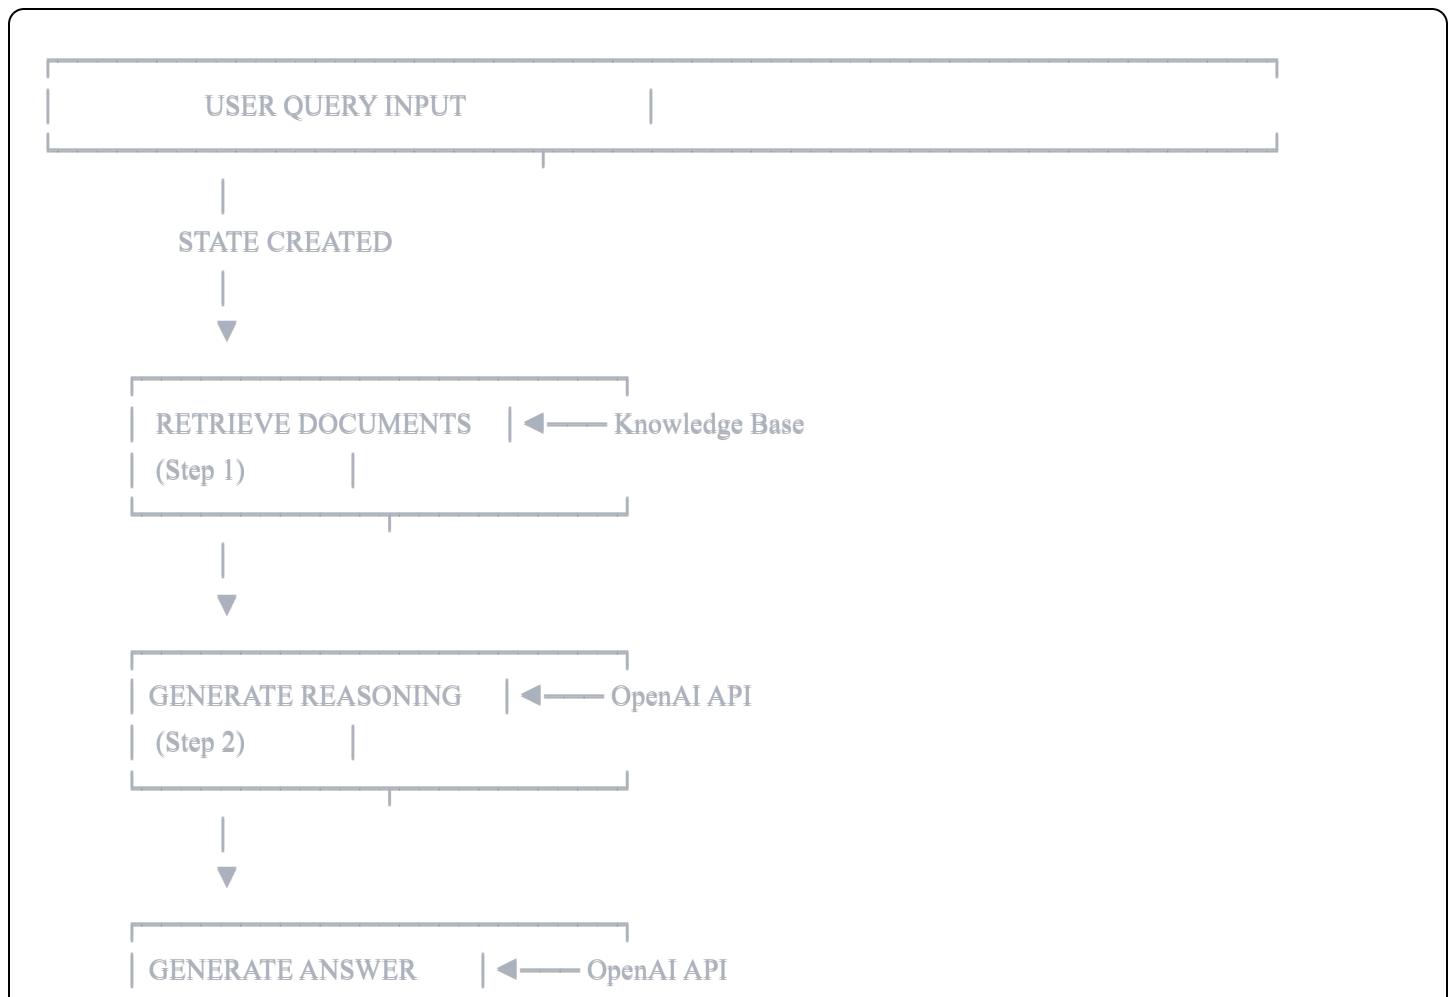
Before starting, ensure you have:

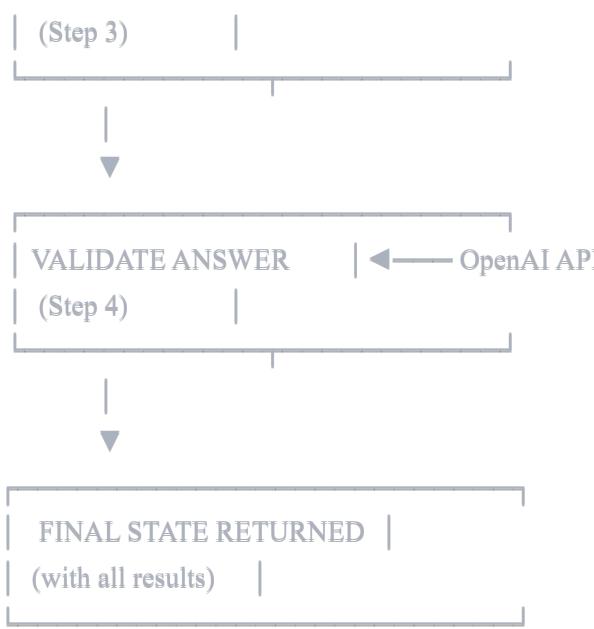
1. **Python 3.9+** installed
2. **OpenAI API Key** (get from <https://platform.openai.com/api-keys>)
3. **Basic Python knowledge** (functions, dictionaries, type hints)
4. **Understanding of LLMs** (what models are and how they work)

## Required Python Version Check

```
bash  
  
python --version # Should be 3.9 or higher
```

## Architecture Overview





## 🎯 Core Concepts

### 1. State Management

**What is State?** State is the data structure that flows through your agent. It carries information between nodes.

```

python

class AgentState(TypedDict):
    input: str                  # User's question
    steps: Annotated[list, add]  # Accumulated workflow steps
    documents: list             # Retrieved documents
    reasoning: str              # LLM reasoning
    final_answer: str           # Final output

```

#### Why TypedDict?

- Type safety at runtime
- IDE autocompletion
- Clear contract for state shape
- Easy debugging

### 2. Annotated Reducer Pattern

```

python

```

```
steps: Annotated[list, add]
```

This is advanced Python typing:

- `Annotated` marks metadata
- `[add]` operator concatenates lists
- Each node can append steps without overwriting

**Without this pattern (✗ wrong):**

```
python  
state["steps"] = ["step1"] # Overwrites previous
```

**With this pattern (✓ correct):**

```
python  
state["steps"].append("step1") # Accumulates
```

### 3. Nodes

A node is a function that:

1. Takes current state as input
2. Performs an action
3. Updates and returns state

```
python  
  
def my_node(state: AgentState) -> AgentState:  
    # Do something  
    state["field"] = "value"  
    return state
```

### 4. Edges

Edges define the flow between nodes:

```
python
```

```
workflow.add_edge("node_a", "node_b")      # Sequential  
workflow.add_conditional_edges(           # Conditional  
    "node_x",  
    decision_function,  
    {"path1": "node_y", "path2": "node_z"}  
)
```

## 5. Graph Compilation

```
python  
graph = workflow.compile()
```

This converts your graph definition into an executable object. It optimizes and validates the workflow.

---

## Code Breakdown

### Section 1: Imports & Dependencies

```
python  
  
from langgraph.graph import StateGraph, START, END  
from typing import TypedDict, Annotated  
from openai import OpenAI  
from operator import add
```

Import	Purpose
StateGraph	Create workflow graph
START, END	Special nodes for entry/exit
TypedDict	Type-safe state definition
Annotated	Add metadata to types
OpenAI	Call OpenAI models
add	Concatenate lists in state

## Section 2: State Definition

```
python

class AgentState(TypedDict):
    input: str          # Raw user input
    steps: Annotated[list, add]  # Workflow progress tracker
    documents: list      # Retrieved context
    reasoning: str        # Intermediate thinking
    final_answer: str     # Final output
```

### Each field explained:

Field	Type	Purpose	Initial Value
input	str	User's question	Passed in
steps	Annotated list	Track executed steps	Empty list
documents	list	Retrieved knowledge	Empty list
reasoning	str	LLM thinking process	Empty string
final_answer	str	Answer to user	Empty string

## Section 3: Initialize OpenAI Client

```
python

client = OpenAI() # Reads OPENAI_API_KEY from environment
```

### How it works:

- Automatically reads `OPENAI_API_KEY` environment variable
- Creates authenticated client
- Ready to make API calls

### Setting the API key:

```
bash
```

```
export OPENAI_API_KEY="sk-proj-..." # Linux/Mac
set OPENAI_API_KEY=sk-proj-...      # Windows
```

## Section 4: Knowledge Base

```
python

KNOWLEDGE_BASE = {
    "python": ["Python is...", "It uses..."],
    "api": ["API stands for...", "REST APIs..."],
    "database": ["Databases store...", "SQL is..."]
}
```

### In production:

- Replace with vector database (Pinecone, Weaviate)
- Connect to SQL/NoSQL database
- Call external APIs
- Load from files

## Section 5: Node Functions

### Node 1: Retrieve Documents

```
python

def retrieve_documents(state: AgentState) -> AgentState:
    """Step 1: Retrieve relevant documents"""

    query = state["input"].lower()
    docs = []

    # Keyword matching (in production: use embeddings + similarity search)
    for topic, content in KNOWLEDGE_BASE.items():
        if topic in query:
            docs.extend(content)

    state["steps"].append("✓ Retrieved documents")
    state["documents"] = docs

    return state
```

## What it does:

1. Extracts user query
2. Searches knowledge base
3. Updates steps tracker
4. Stores documents in state
5. Returns updated state

## Real-world improvement:

```
python

# Use embeddings for better matching
from openai import OpenAI
client = OpenAI()

query_embedding = client.embeddings.create(
    model="text-embedding-3-small",
    input=state["input"]
)
# Compare with document embeddings
```

## Node 2: Generate Reasoning

```
python
```

```

def generate_reasoning(state: AgentState) -> AgentState:
    """Step 2: Generate reasoning using OpenAI"""

    documents_text = "\n".join(state["documents"])

    response = client.messages.create(
        model="gpt-4o-mini",
        max_tokens=200,
        messages=[{
            "role": "user",
            "content": f"""Explain how these docs relate to: {state['input']}"""
        }
    )

    state["reasoning"] = response.content[0].text
    state["steps"].append("✓ Generated reasoning")

    return state

```

## Key concepts:

- `messages.create()` - Call OpenAI API
- `model="gpt-4o-mini"` - Fast, cheap model
- `max_tokens=200` - Limit response length
- Prompt engineering included
- Chain-of-thought reasoning

## Node 3: Generate Answer

python

```
def generate_answer(state: AgentState) -> AgentState:
```

```
    """Step 3: Generate final answer"""


```

```
    response = client.messages.create(
```

```
        model="gpt-4o-mini",
```

```
        max_tokens=300,
```

```
        messages=[{
```

```
            "role": "user",
```

```
            "content": f"""Answer this question:
```

```
Question: {state['input']}
```

```
Knowledge: {"\n".join(state["documents"])}  
Reasoning: {state["reasoning"]}
```

```
Provide a clear answer."""
    ]
```

```
)
```

```
state["final_answer"] = response.content[0].text
```

```
state["steps"].append("✓ Generated final answer")
```

```
return state
```

## Why multiple steps?

- Reasoning helps LLM think through problem
- Better answers than direct approach
- Follows chain-of-thought pattern
- More transparent process

## Node 4: Validate Answer

```
python
```

```

def validate_answer(state: AgentState) -> AgentState:
    """Step 4: Validate answer quality"""

    response = client.messages.create(
        model="gpt-4o-mini",
        max_tokens=100,
        messages=[{
            "role": "user",
            "content": f"""Is this answer helpful? YES or NO:
Question: {state['input']}
Answer: {state['final_answer']}"""
        }]
    )

    validation = response.content[0].text.strip()
    state["steps"].append(f"✓ Validated answer ({validation})")

    return state

```

### Real-world use:

- Filter poor answers
- Trigger re-generation if NO
- Add confidence scoring
- Log for quality monitoring

### Section 6: Build the Graph

python

```

workflow = StateGraph(AgentState)

# Add all nodes
workflow.add_node("retrieve", retrieve_documents)
workflow.add_node("reasoning", generate_reasoning)
workflow.add_node("answer", generate_answer)
workflow.add_node("validate", validate_answer)

# Connect nodes sequentially
workflow.add_edge(START, "retrieve")      # Entry point
workflow.add_edge("retrieve", "reasoning") # Step 1 → 2
workflow.add_edge("reasoning", "answer")   # Step 2 → 3
workflow.add_edge("answer", "validate")    # Step 3 → 4
workflow.add_edge("validate", END)        # Exit point

# Compile to executable
graph = workflow.compile()

```

## Graph visualization:

START → retrieve → reasoning → answer → validate → END

## Execution Flow

### Step-by-step execution:

#### 1. User Input

```

python

query = "What is Python and why is it useful?"

```

#### 2. Initialize State

```
python
```

```
initial_state = AgentState(  
    input="What is Python and why is it useful?",  
    steps=[],  
    documents=[],  
    reasoning="",  
    final_answer=""  
)
```

### 3. Invoke Graph

```
python  
final_state = graph.invoke(initial_state)
```

### 4. Node 1: Retrieve

- Search KB for "python"
- Find 3 documents
- state["documents"] = ["Python is...", "It uses...", ...]
- state["steps"] = ["✓ Retrieved documents"]

### 5. Node 2: Reasoning

- Call OpenAI with documents + question
- Generate: "Python is relevant because..."
- state["reasoning"] = "Python is relevant..."
- state["steps"] = [..., "✓ Generated reasoning"]

### 6. Node 3: Answer

- Call OpenAI with all context
- Generate full answer
- state["final\_answer"] = "Python is a high-level language..."
- state["steps"] = [..., "✓ Generated final answer"]

### 7. Node 4: Validate

- Ask OpenAI: "Is this good?"
- Get: "YES"
- state["steps"] = [..., "✓ Validated answer (YES)"]

## 8. Return Final State

```
python

{
  "input": "What is Python...",
  "steps": ["✓ Retrieved documents", "✓ Generated reasoning", ...],
  "documents": ["Python is...", ...],
  "reasoning": "Python is relevant...",
  "final_answer": "Python is a high-level language..."
}
```

## 🚀 Installation & Setup

### Step 1: Create Virtual Environment

```
bash

# Create
python -m venv venv

# Activate
source venv/bin/activate      # Linux/Mac
venv\Scripts\activate         # Windows
```

### Step 2: Install Dependencies

```
bash

pip install langgraph openai python-dotenv
```

### What each package does:

Package	Purpose
langgraph	Workflow orchestration
openai	OpenAI API client
python-dotenv	Load environment variables

## Step 3: Set API Key

### Option A: Environment Variable (Recommended)

```
bash
```

```
export OPENAI_API_KEY="sk-proj-your-key-here"
```

### Option B: .env File

Create `.env`:

```
OPENAI_API_KEY=sk-proj-your-key-here
```

Load in code:

```
python
```

```
from dotenv import load_dotenv
import os

load_dotenv()
api_key = os.getenv("OPENAI_API_KEY")
client = OpenAI(api_key=api_key)
```

## Step 4: Verify Installation

```
bash
```

```
python -c "import langgraph; import openai; print('✓ All packages installed')"
```

## 📝 Usage Examples

### Example 1: Basic Usage

```
python
```

```
from_agent_import run_agent
```

```
result = run_agent("What is Python?")
print(result["final_answer"])
```

## Output:

Python is a high-level, interpreted programming language known for its simplicity and readability. It uses indentation for code blocks and has powerful frameworks like Django and Flask for web development.

## Example 2: Multiple Queries

```
python

queries = [
    "What is Python and why is it useful?",
    "Explain REST APIs to me",
    "Tell me about databases"
]

for query in queries:
    print(f"\n ? Query: {query}")
    result = run_agent(query)
    print(f" ✓ Answer: {result['final_answer'][:100]}")
```

## Example 3: Custom Knowledge Base

```
python

KNOWLEDGE_BASE = {
    "machine-learning": [
        "ML is subset of AI",
        "Common algorithms: regression, classification",
        "Libraries: TensorFlow, PyTorch, scikit-learn"
    ],
    "deep-learning": [
        "Deep learning uses neural networks",
        "Popular architectures: CNN, RNN, Transformers",
        "Used in NLP and computer vision"
    ]
}

result = run_agent("How does deep learning work?")
```

## Example 4: Extract Results

```
python
```

```
final_state = graph.invoke(initial_state)

# Get specific outputs
answer = final_state["final_answer"]
reasoning = final_state["reasoning"]
steps = final_state["steps"]
docs = final_state["documents"]

print("==== WORKFLOW EXECUTION ====")
for step in steps:
    print(step)

print("\n==== DOCUMENTS USED ====")
for doc in docs:
    print(f"• {doc}")

print("\n==== REASONING ====")
print(reasoning)

print("\n==== FINAL ANSWER ====")
print(answer)
```

## 🔧 Advanced Customization

### 1. Add Conditional Routing

python

```
def should_revalidate(state: AgentState) -> str:  
    """Decide if we need to regenerate answer"""  
    if "NO" in state["steps"][-1]:  
        return "regenerate"  
    return "end"  
  
workflow.add_conditional_edges(  
    "validate",  
    should_revalidate,  
    {  
        "regenerate": "answer",  
        "end": END  
    }  
)
```

## 2. Add Parallel Execution

```
python  
  
from langgraph.graph import StateGraph  
  
# Create branches that execute in parallel  
workflow.add_node("branch_a", func_a)  
workflow.add_node("branch_b", func_b)  
  
workflow.add_edge("start", "branch_a")  
workflow.add_edge("start", "branch_b")  
  
# Both branches merge  
workflow.add_edge(["branch_a", "branch_b"], "merge_node")
```

## 3. Add Error Handling

```
python
```

```

def retrieve_documents_safe(state: AgentState) -> AgentState:
    """Retrieve with error handling"""

    try:
        query = state["input"].lower()
        docs = []
        for topic, content in KNOWLEDGE_BASE.items():
            if topic in query:
                docs.extend(content)
        state["documents"] = docs
    except Exception as e:
        state["documents"] = [f"Error: {str(e)}"]

    state["steps"].append("✓ Retrieved documents")
    return state

```

## 4. Add Logging

```

python

import logging

logging.basicConfig(level=logging.INFO)
logger = logging.getLogger(__name__)

def retrieve_documents(state: AgentState) -> AgentState:
    """Retrieve with logging"""

    logger.info(f'Retrieving for query: {state["input"]}')

    # ... retrieval logic ...

    logger.info(f'Found {len(state["documents"])} documents')
    return state

```

## 5. Custom Model Selection

```

python

```

```

def generate_reasoning(state: AgentState) -> AgentState:
    """Use different models for different cases"""

    if len(state["input"]) < 50:
        model = "gpt-4o-mini" # Simple queries
    else:
        model = "gpt-4"      # Complex queries

    response = client.messages.create(
        model=model,
        max_tokens=200,
        messages=[...]
    )

    return state

```

## 6. Add Metrics & Monitoring

```

python

import time

def retrieve_documents(state: AgentState) -> AgentState:
    """Retrieve with metrics"""
    start = time.time()

    # ... retrieval logic ...

    duration = time.time() - start
    state["steps"].append(f"✓ Retrieved documents ({duration:.2f}s)")

    return state

```

## Troubleshooting

### Problem 1: API Key Not Found

Error:

AuthenticationError: Error code: 401

## Solution:

```
bash

# Verify key is set
echo $OPENAI_API_KEY

# If empty, set it
export OPENAI_API_KEY="sk-proj-..."

# Or create .env file
echo "OPENAI_API_KEY=sk-proj-..." > .env
```

## Problem 2: Rate Limiting

### Wrong:

```
RateLimitError: Error code: 429
```

### Solution:

```
python

import time

def retry_with_backoff(func, max_retries=3):
    for attempt in range(max_retries):
        try:
            return func()
        except Exception as e:
            if attempt < max_retries - 1:
                wait_time = 2 ** attempt
                print(f"Rate limited. Waiting {wait_time}s...")
                time.sleep(wait_time)
            else:
                raise
```

## Problem 3: State Not Updating

### Wrong:

```
python
```

```
state["documents"] = [] # This might not persist
```

**Correct:**

```
python

state["documents"] = retrieved_docs # Ensure full assignment
return state # Always return state
```

## Problem 4: Nodes Not Executing

**Check:**

```
python

# Verify edges are connected
print(graph.get_graph().nodes)
print(graph.get_graph().edges)

# Visualize
graph.get_graph().draw_mermaid()
```

## Problem 5: High API Costs

**Solution:**

```
python

# Use cheaper model
model = "gpt-4o-mini" # 5x cheaper than gpt-4

# Reduce tokens
max_tokens = 100 # Instead of 1000

# Cache responses
from functools import lru_cache

@lru_cache(maxsize=128)
def get_answer(query: str):
    # Only call API once per unique query
    pass
```

# Performance Optimization

## 1. Batch Processing

```
python

def run_batch(queries: list) -> list:
    """Process multiple queries efficiently"""
    results = []
    for query in queries:
        initial_state = AgentState(
            input=query,
            steps=[],
            documents=[],
            reasoning="",
            final_answer=""
        )
        result = graph.invoke(initial_state)
        results.append(result)
    return results
```

## 2. Caching

```
python

from functools import lru_cache

@lru_cache(maxsize=256)
def cached_retrieve(query: str) -> tuple:
    """Cache document retrieval"""
    docs = []
    for topic, content in KNOWLEDGE_BASE.items():
        if topic in query.lower():
            docs.extend(content)
    return tuple(docs)
```

## 3. Async Execution

```
python
```

```
import asyncio

async def run_agent_async(query: str):
    """Non-blocking execution"""
    loop = asyncio.get_event_loop()
    result = await loop.run_in_executor(None, run_agent, query)
    return result

# Usage
results = asyncio.gather(
    run_agent_async("Query 1"),
    run_agent_async("Query 2"),
    run_agent_async("Query 3")
)
```

---

## 🎓 Learning Path

1. **Understand State** - How data flows through nodes
  2. **Build Simple Graph** - 2-3 nodes sequentially
  3. **Add Conditional Logic** - Route based on decisions
  4. **Integrate Tools** - Call external APIs/databases
  5. **Add Error Handling** - Handle failures gracefully
  6. **Optimize Performance** - Cache, batch, async
  7. **Deploy to Production** - Docker, serverless, cloud
- 

## 📚 Additional Resources

- **LangGraph Docs:** <https://langchain-ai.github.io/langgraph/>
  - **OpenAI API Docs:** <https://platform.openai.com/docs>
  - **Chain-of-Thought Papers:** <https://arxiv.org/abs/2201.11903>
  - **Agentic AI Patterns:** <https://github.com/langchain-ai/langgraph>
-

## Checklist for Production

- API keys securely stored (use .env)
  - Error handling for all nodes
  - Logging at each step
  - Rate limiting implemented
  - Cost monitoring enabled
  - Tests for each node function
  - Documentation for team
  - Performance benchmarks
  - Security audit completed
  - Monitoring/alerting setup
- 

Happy building! 